

# Introduction to Operating Systems and SQL for Data Science

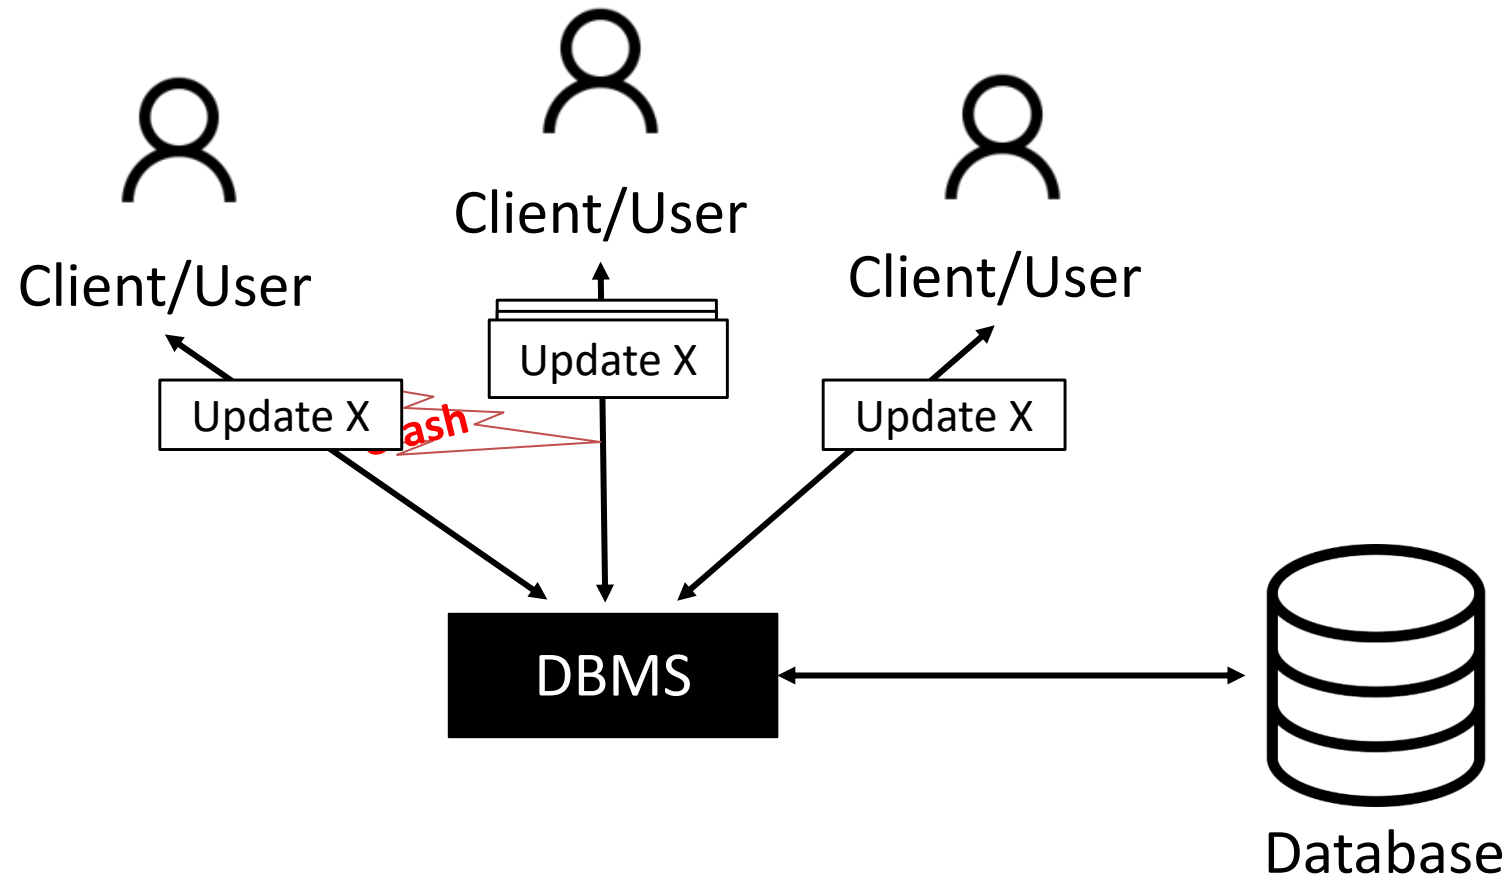
---

## Lecture 12 – Transactions

# Previous lecture

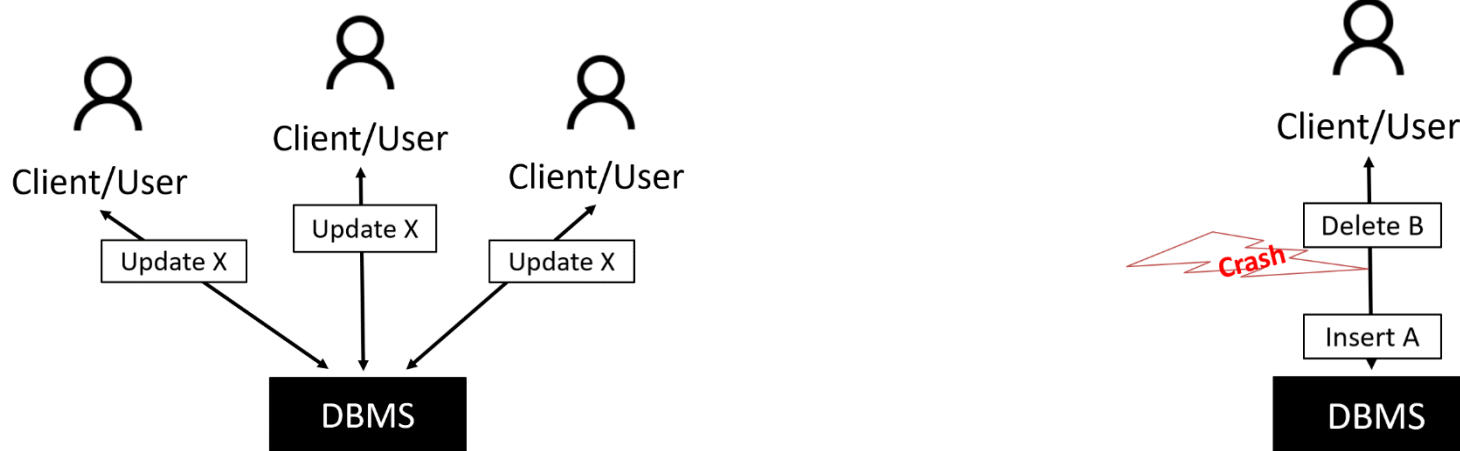
- ERD
- Entity
- attributes
- Connection between entities
- Connection types
- generalization

# Challenges in a real database



# Challenges in a real database

Challenges in a DB with **many users** that may **fail**



**Transactions to save the day!**

# Lecture topics

- Problems in a multi-user DB that can fail
- **What is a transaction?**
  - What should a DBMS do with transactions?
- How DBMSs **implement transactions?**
- Trading data **consistency for speed**

# Transactions

# Example #1: Student Cleanup

## Students

ID	Name	GPA	Status
1	Avi	82.5	Active
2	Benny	62.3	Rejected
3	Danna	91.2	Active
4	Ehud	65.3	Rejected

## Archive

ID	Name	GPA	Status
----	------	-----	--------

Is it possible for a student to be in both tables?  
(In a normal condition - **no!**)

Insert into **Archive**

Select \* from **Students** Where **Status** = 'Rejected'

Delete From **Students** Where **Status** = 'Rejected'

# Example #1: Student Cleanup – normal state

## Students

ID	Name	GPA	Status
1	Avi	82.5	Active
3	Danna	91.2	Active

## Archive

ID	Name	GPA	Status
2	Benny	62.3	Rejected
4	Ehud	65.3	Rejected

Insert into **Archive**

Select \* from **Students** Where **Status** = 'Rejected'

Delete From **Students** Where **Status** = 'Rejected'



# Example #1: Student Cleanup

## Students

ID	Name	GPA	Status
1	Avi	82.5	Active
3	Danna	91.2	Active

## Archive

ID	Name	GPA	Status
2	Benny	62.3	Rejected
4	Ehud	65.3	Rejected

**What if the DBMS crashed after the INSERT?**

Insert into **Archive**

Select \* from **Students** Where **Status** = 'Rejected'

Delete From **Students** Where **Status** = 'Rejected'

# Example #1: Student Cleanup

## Students

ID	Name	GPA	Status
1	Avi	82.5	Active
2	Benny	62.3	Rejected
3	Danna	91.2	Active
4	Ehud	65.3	Rejected

## Archive

ID	Name	GPA	Status
----	------	-----	--------

Insert into **Archive**

Select \* from **Students** Where **Status** = 'Rejected'

Delete From **Students** Where **Status** = 'Rejected'

# Example #1: Student Cleanup

## Students

ID	Name	GPA	Status
1	Avi	82.5	Active
2	Benny	62.3	Rejected
3	Danna	91.2	Active
4	Ehud	65.3	Rejected

## Archive

ID	Name	GPA	Status
2	Benny	62.3	Rejected
4	Ehud	65.3	Rejected



Crash

Insert into **Archive**

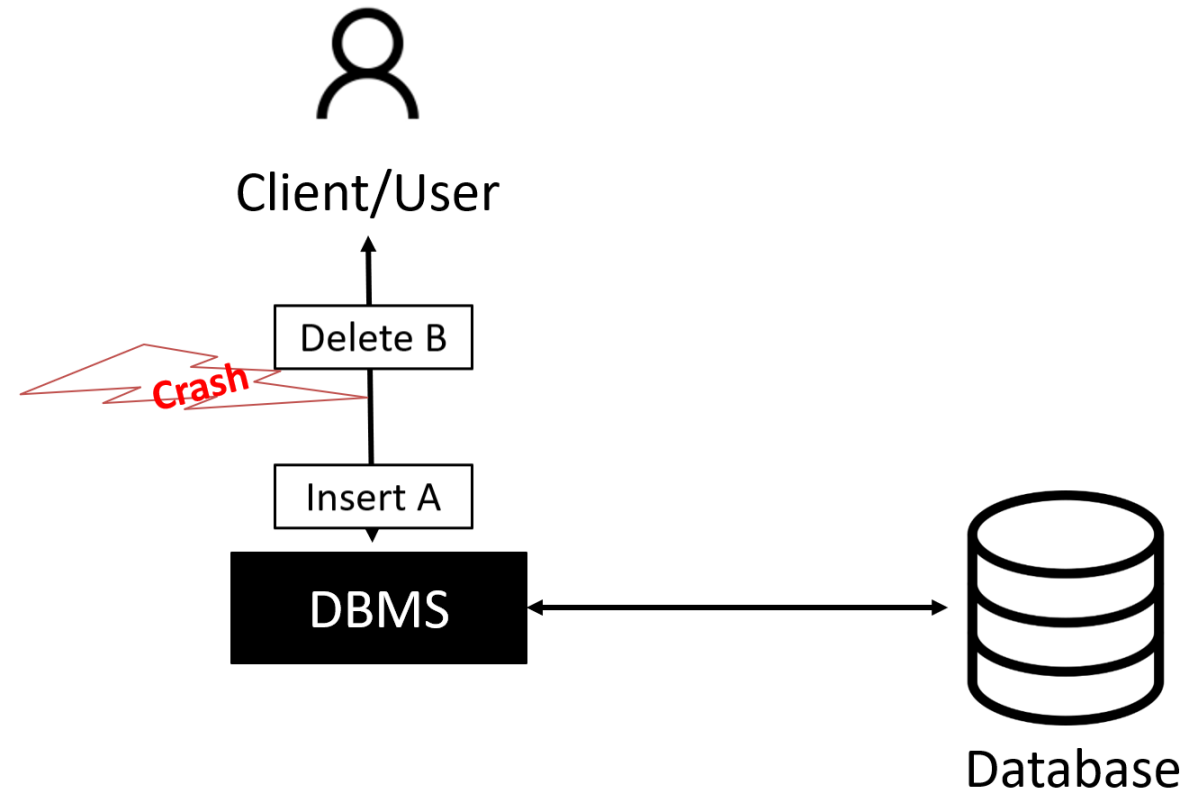
Select \* from **Students** Where **Status** = 'Rejected'

Delete From **Students** Where **Status** = 'Rejected'

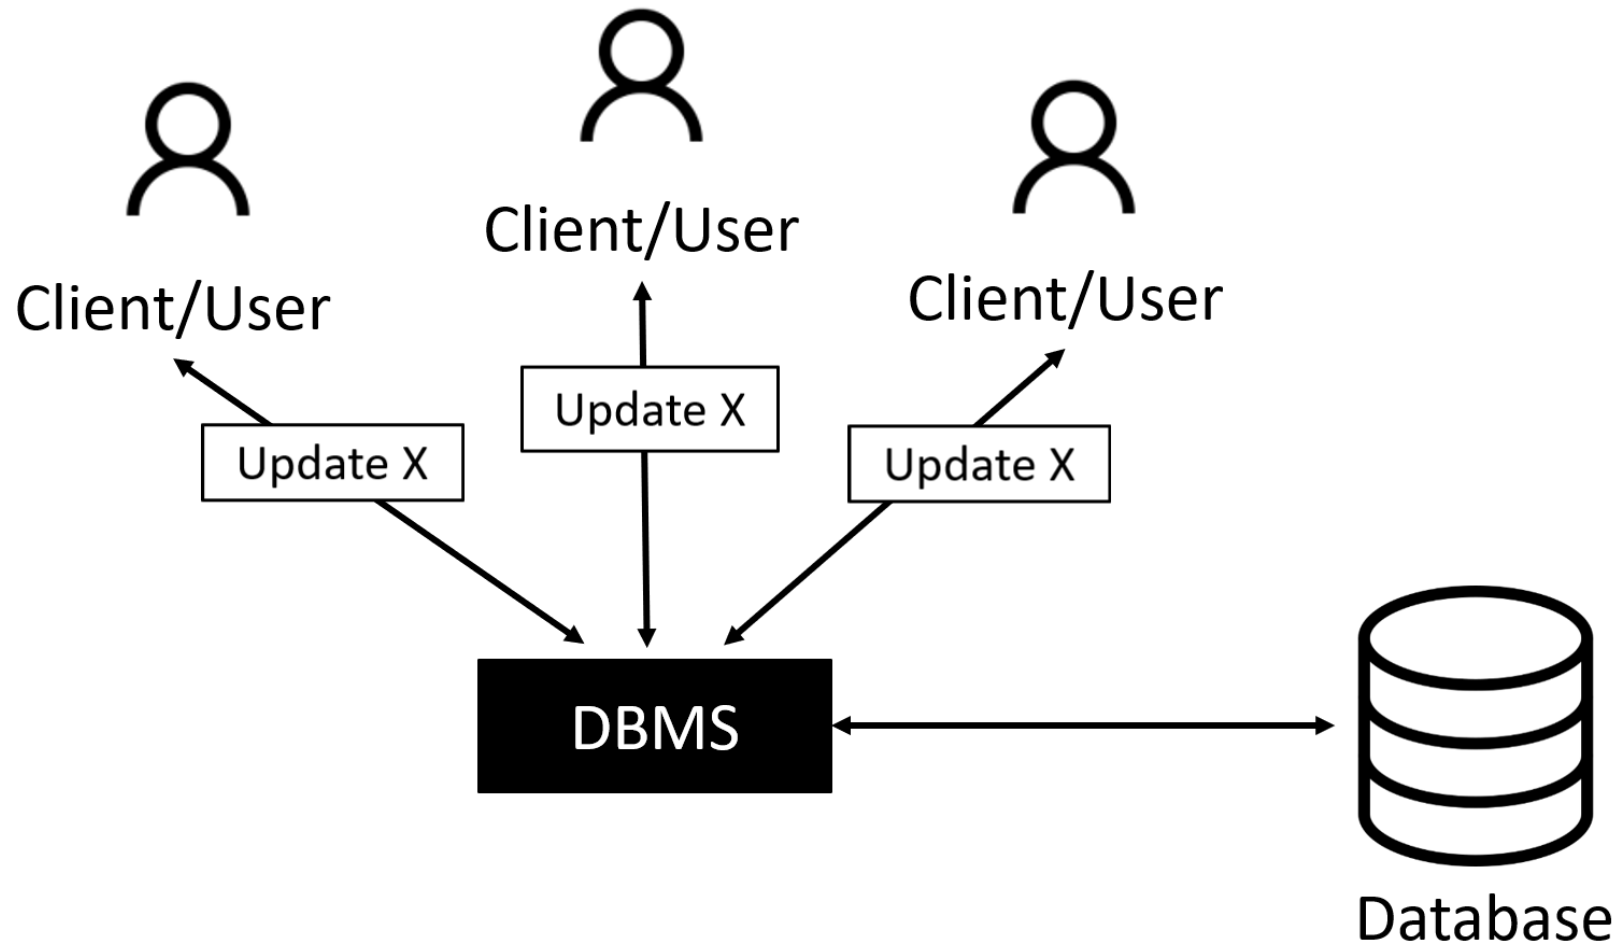
# Example #1: Student Cleanup

- There will be duplicate records both in Archive and Students tables
- If, we run again the query we get failure (duplicate records with the same primary key)
- We need a way to say execute **“all or nothing”**

# Example #1: DBMS crash



# Example #2: Multiple Users



# How to perform an update

Update **Employees**

Set **salary = salary + 1000**

Where **SSN= 7**

# Problem #1: Concurrent Updates

Update **Employees**

Set **salary = salary + 1000**

Where **SSN= 7**

Update **Employees**

Set **salary = salary + 1500**

Where **SSN= 7**

**Concurrently**

Is it 14,500?

Maybe 13,500 or  
13,000???

Employees

SSN	Name	Salary
7	Danna	12,000

**What is Danna's salary?**



# Problem #2: Concurrent Updates

Update **Employees**

Set **salary = salary + 2000**

Where **SSN= 7**

Update **Employees**

Set **salary = salary/2**

Where **SSN= 7**

**Concurrently**

Is it 7,000 or 8,000? (even  
worse than the previous  
example – no final stage)

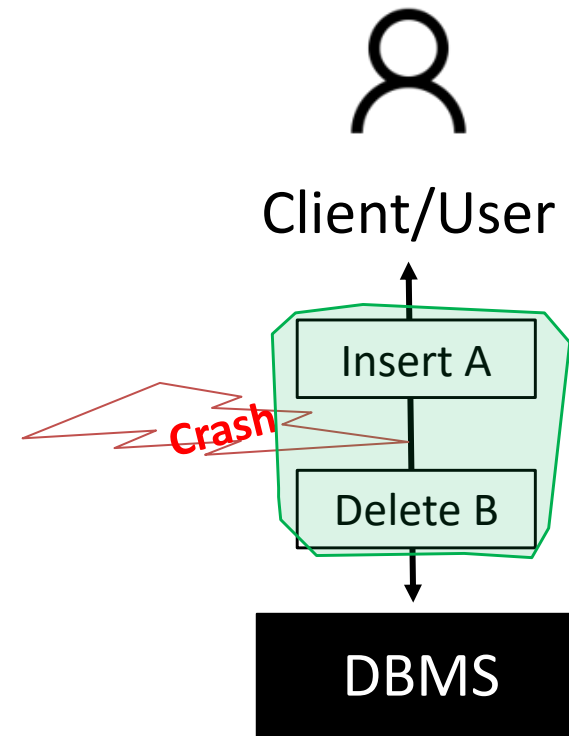
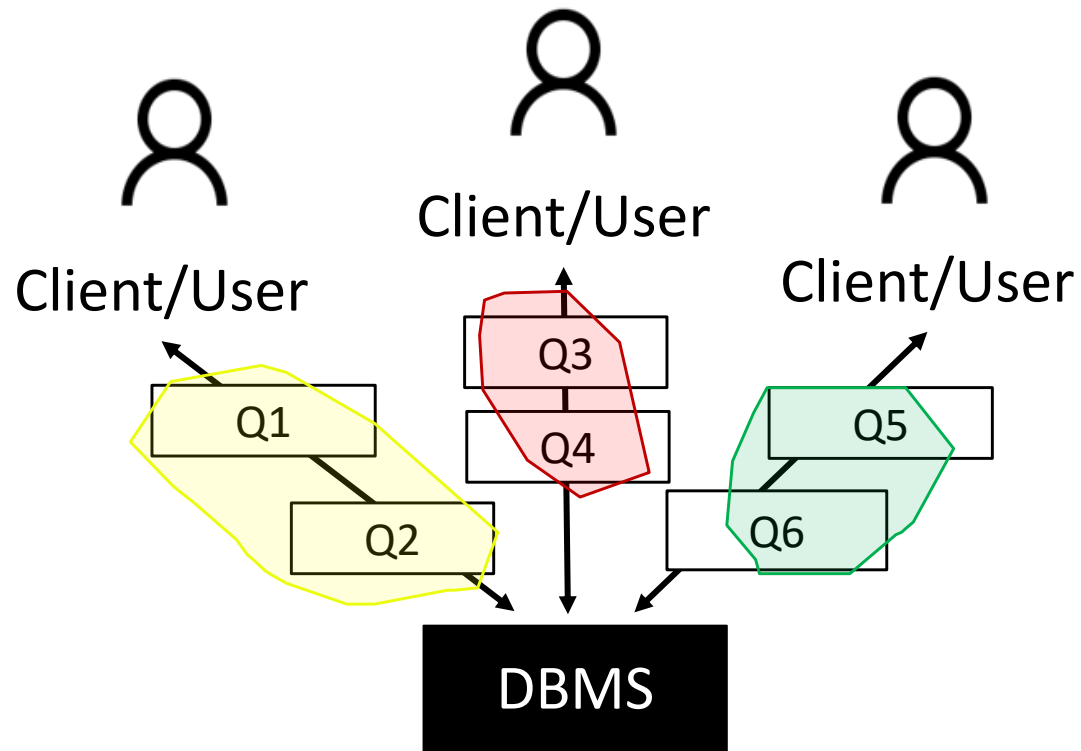
Maybe 14,000 or  
6,000???

Employees

SSN	Name	Salary
7	Danna	12,000

**What is Danna's salary?**

# Actions Need to be Grouped Together!

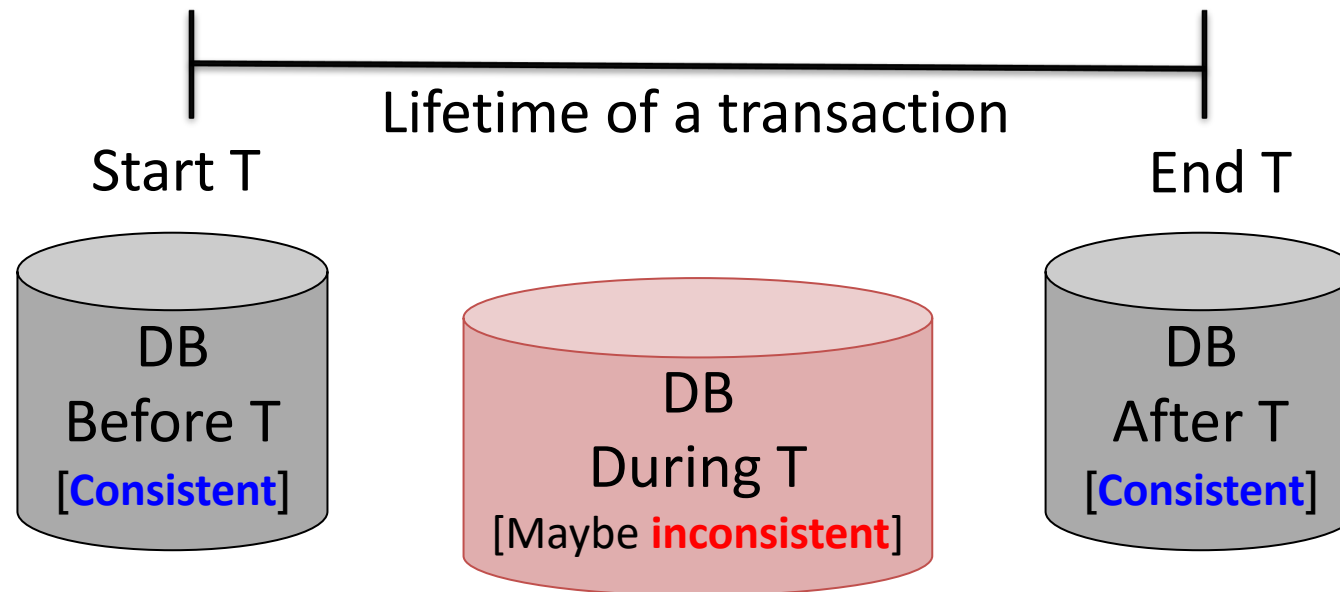


# Transactions

# Transactions to the Rescue!

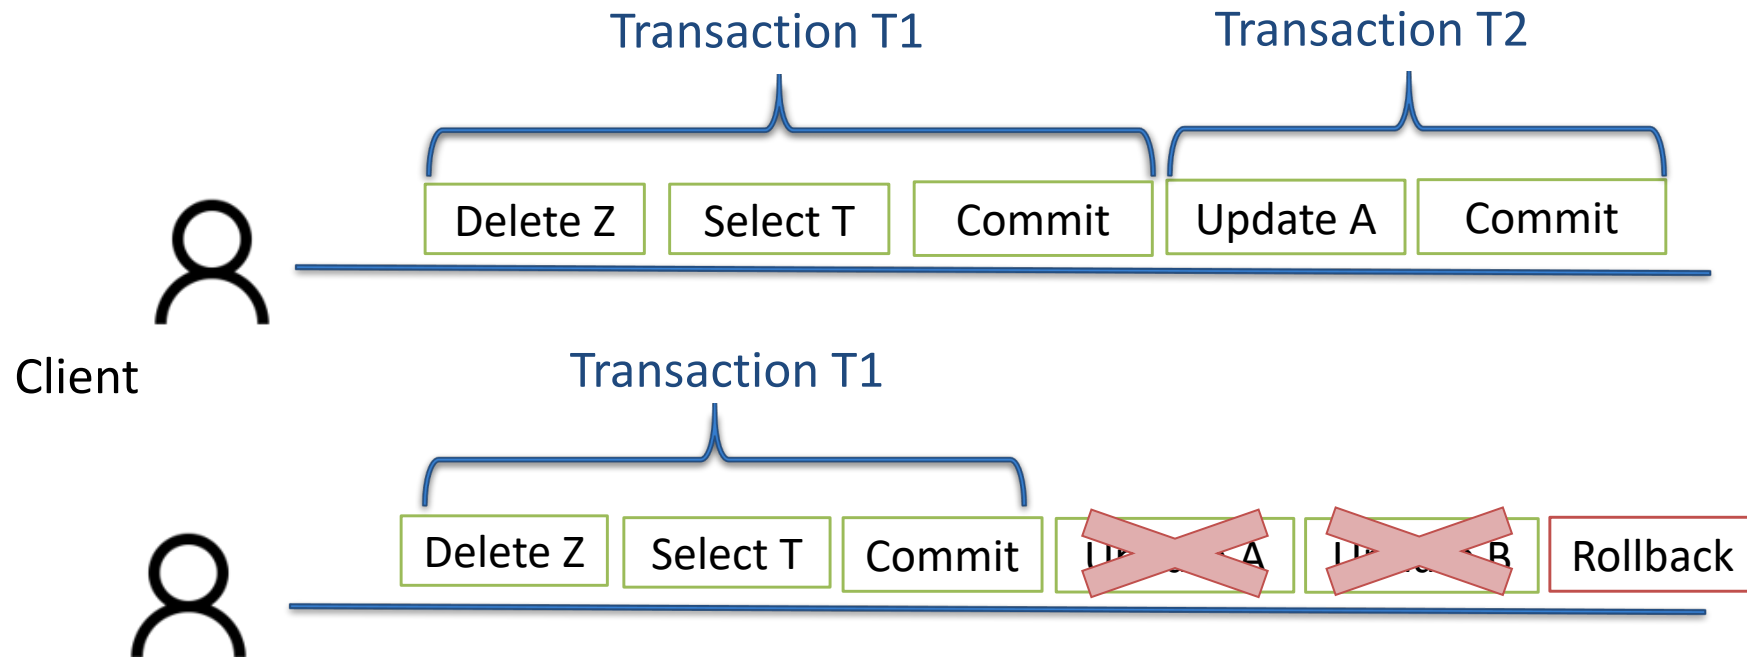
**Transaction:** a sequence of SQL operations that are considered as a single unit

- In a transaction either **all actions are done or none**
- Transitions the DB from one consistent state to another

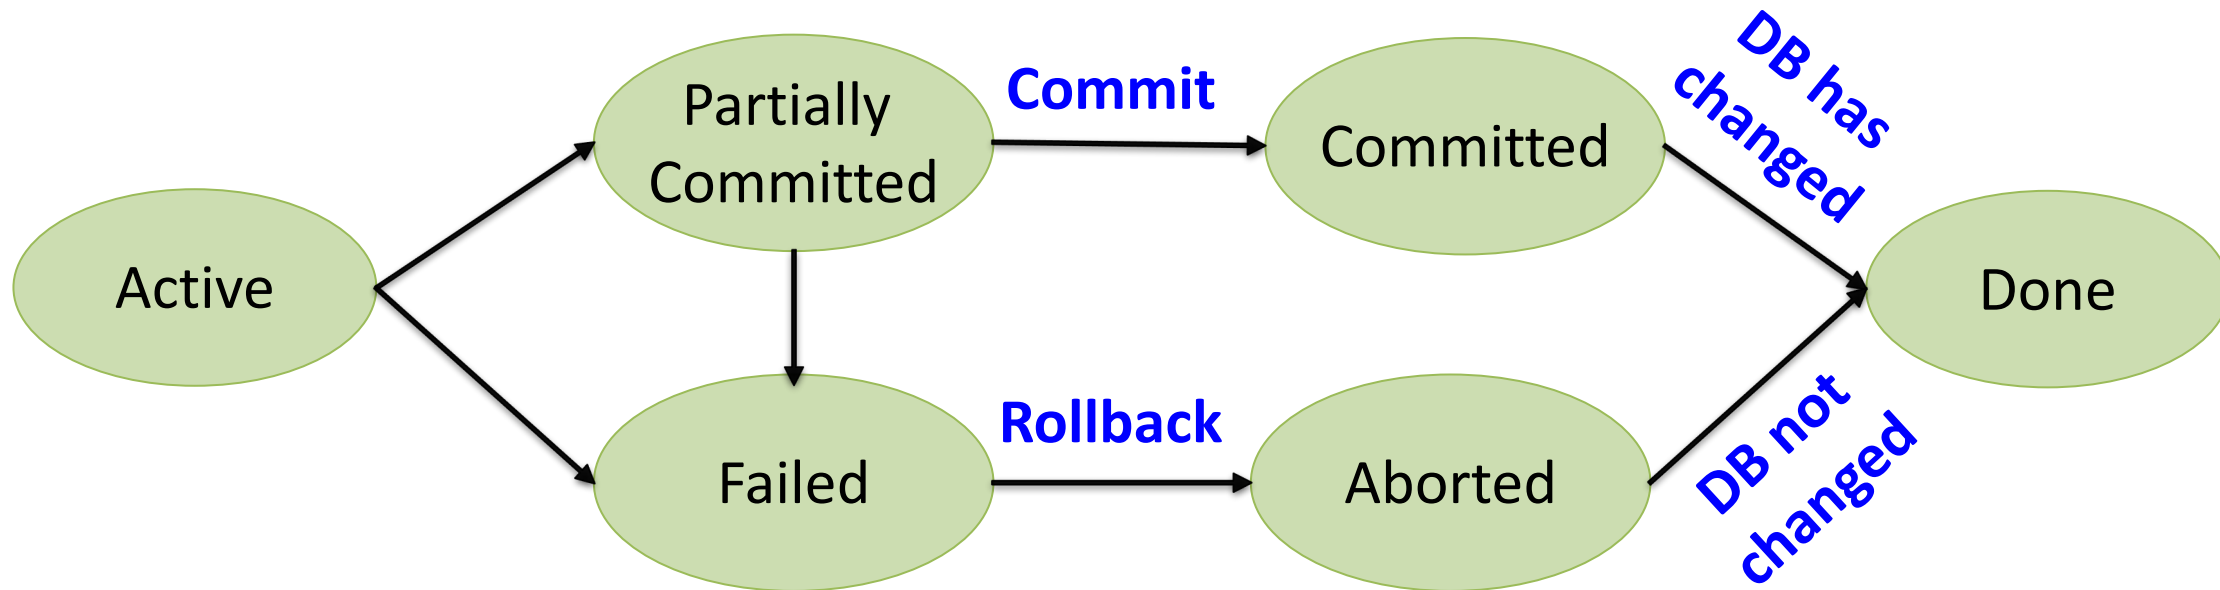


# Transaction mechanism

- How we create transactions?
  1. We apply SQL actions.
  2. When we want to “close” a Transaction we call this command **Commit**.
  3. If we want to “cancel” a started Transaction we do a **Rollback**.
- In many languages we can define auto commit



# Transaction lifecycle



ACID

# ACID: Desired Properties of Transactions

- Atomicity
- Consistency
- Isolation
- Durability



# ACID: Atomicity

- Each transaction is “all-or-nothing” never left half done.
- What to do if an error occurs in the middle of a Transaction?
  - We perform a rollback to the Transaction to "clear" the changes

# Atomicity

## Students

ID	Name	GPA	Status
1	Avi	82.5	Active
2	Benny	62.3	Rejected
3	Danna	91.2	Active
4	Ehud	65.3	Rejected

## Archive

ID	Name	GPA	Status
2	Benny	62.3	Rejected
4	Ehud	65.3	Rejected



Crash

We solve this problem with atomicity as we return to the commit stage.

Insert into **Archive**

Select \* from **Students** Where **Status** = 'Rejected'

Delete From **Students** Where **Status** = 'Rejected'

# ACID: Consistency

Each client, each transaction:

- If all **constraints hold when transaction begins**
- Guarantee all **constraints hold when transaction ends**

- What kind of database constraints are there?

Referential integrity(e.g., cascade delete)

- What is an inconsistent state?

# Atomicity vs. Consistency: are they the same?

- **Atomicity** - In an atomic transaction, a series of database operations either all occur, or nothing occurs. A guarantee of atomicity prevents updates to the database occurring only partially, which can cause greater problems than rejecting the whole series outright.
- **Consistency** - In database systems, a consistent transaction is one that does not violate any integrity constraints during its execution. If a transaction leaves the database in an illegal state, it is aborted, and an error is reported.

# Atomicity vs. Consistency: are they the same?

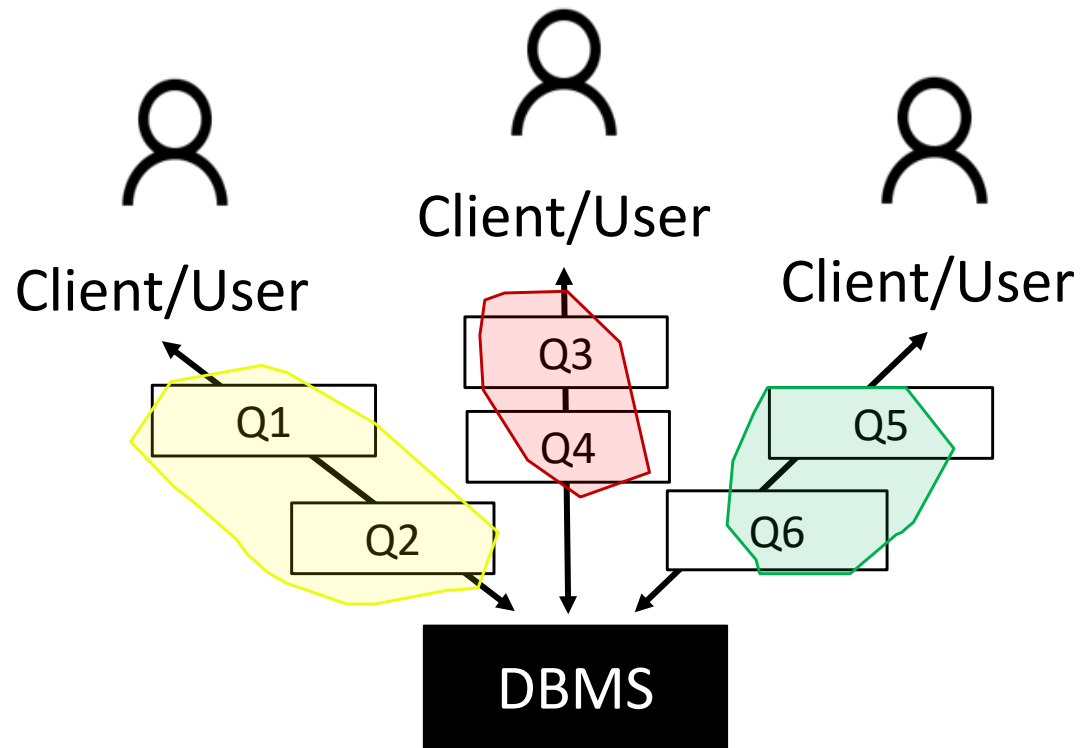
## (Simple words)

- Atomicity – all the transaction or nothing.
- Consistency – guarantee all constraints in the database.  
Such as, Referential integrity (e.g., cascade delete)

# ACID: Isolation

## Isolation

The transaction will behave as if it is the only operation being performed upon the database



# Problem #1: Concurrent Updates

Update **Employees**

Set **salary = salary + 1000**

Where **SSN= 7**

Update **Employees**

Set **salary = salary + 1500**

Where **SSN= 7**

**Concurrently**

Is it 14,500?

Maybe 13,500 or  
13,000???

Employees

SSN	Name	Salary
7	Danna	12,000

**What is Danna's salary?**

# Problem #2: Concurrent Updates

Update **Employees**

Set **salary = salary + 2000**

Where **SSN= 7**

Update **Employees**

Set **salary = salary/2**

Where **SSN= 7**

**Concurrently**

Is it 7,000 or 8,000? (even  
worse than the previous  
example – no final stage)

Maybe 14,000 or  
6,000???

Employees


SSN	Name	Salary
7	Danna	12,000

**What is Danna's salary?**



# ACID: Isolation

- Serializability - Execution must be equivalent to some sequential (serial) order of all transactions
- Often implemented by a locking mechanism



# Problem #2: Concurrent Updates

- What is the right answer?

Update **Employees**  
Set **salary = salary + 2000**  
Where **SSN = 7**

Update **Employees**  
Set **salary = salary/2**  
Where **SSN = 7**

Concurrently

- The thing is there is no one answer...

Employees

SSN	Name	Salary
7	Danna	12,000


What is Danna's salary?

- With Isolation we will get both 7,000 and 8,000 as valid answers but not 14,000 and 6,000.
- As we make sure that transaction preformed in a serialize manner.

# Implementing Transactions on DMBS

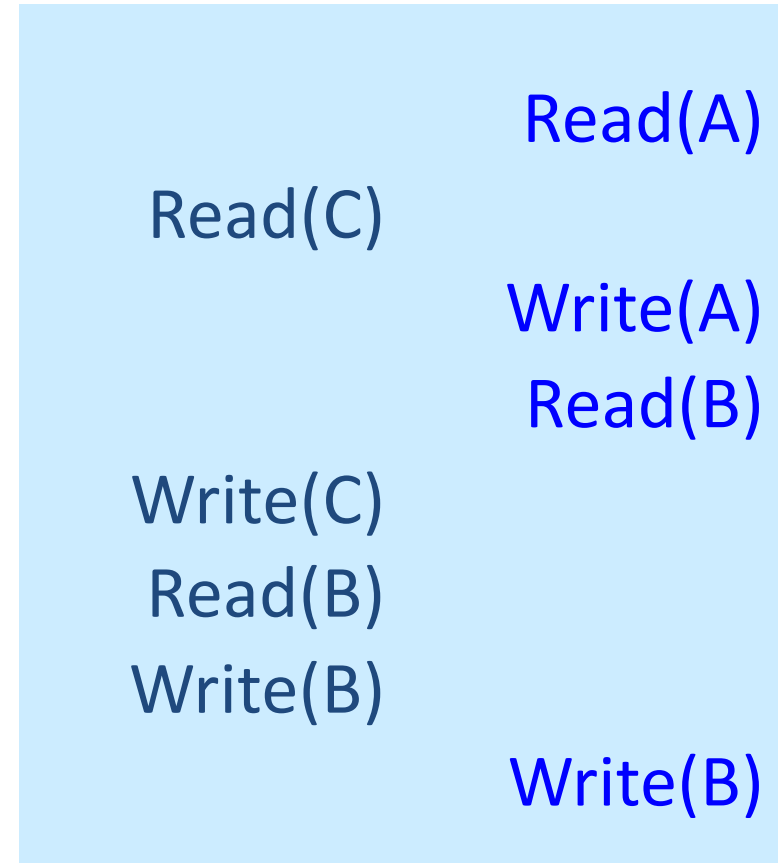
# Goal: maintaining Isolation

How to preserve Isolation? (while executing concurrently)



# Scheduling

- Scheduling is the execution of several Transactions in parallel.
- Only one action is performed at each point in time.
- A reading / writing operation is performed in an atomic form.
- The timing represents the order on the timeline (the vertical axis) in which the actions are performed.



# Serializable scheduling


- Scheduling where transactions are performed one after the other.
- Serial timing is correct.  
(serializable)

Read(A)  
Write(A)  
Read(B)  
Write(B)

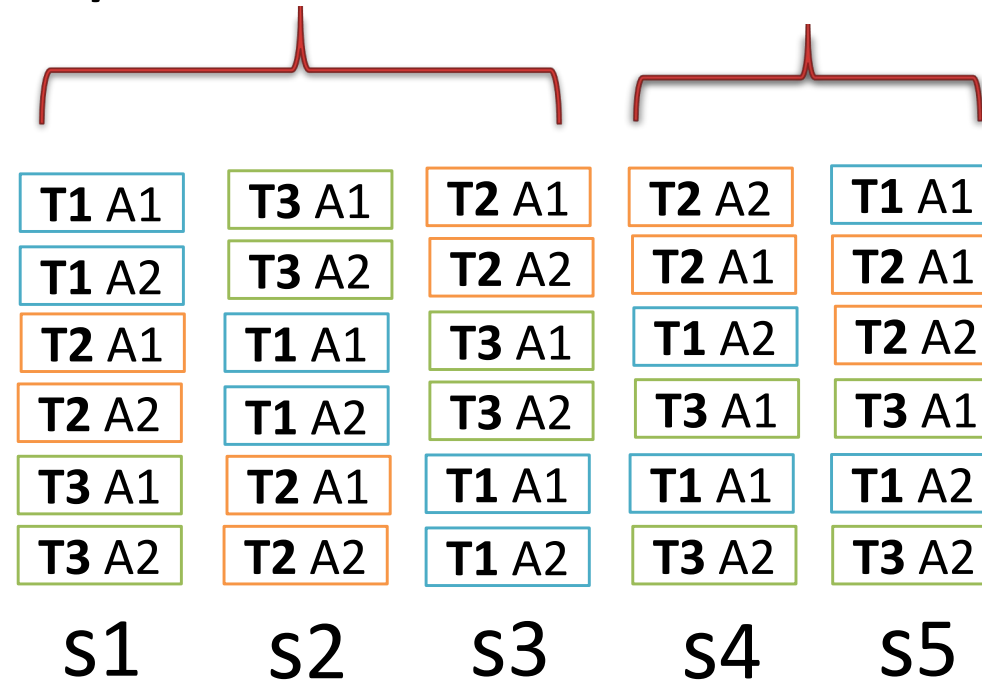
Read(C)  
Write(C)  
Read(B)  
Write(B)

# What is Isolation?

- What is a **correct** schedule?



## Sequential schedules



# Serializable schedule

- A schedule is a serializable schedule if it is equivalent to any serial scheduling.
- Serializable schedule is a correct schedule.
- Need a way to determine that a schedule is serializable without knowing what calculations are being performed.
- That is, for every possible calculation, the timing should be equivalent to a serial scheduling.



# Equivalent schedules

## Two schedules are equivalent if:

- They are composed of the same actions.
- Both schedules have the same effect on the database, meaning that the database has the same values at the end of each of the schedules.
- Both schedules produce the same output to the user, i.e., present the user with the same values for each item they read.

# How the DBMS Achieves Isolation?

## Serializability

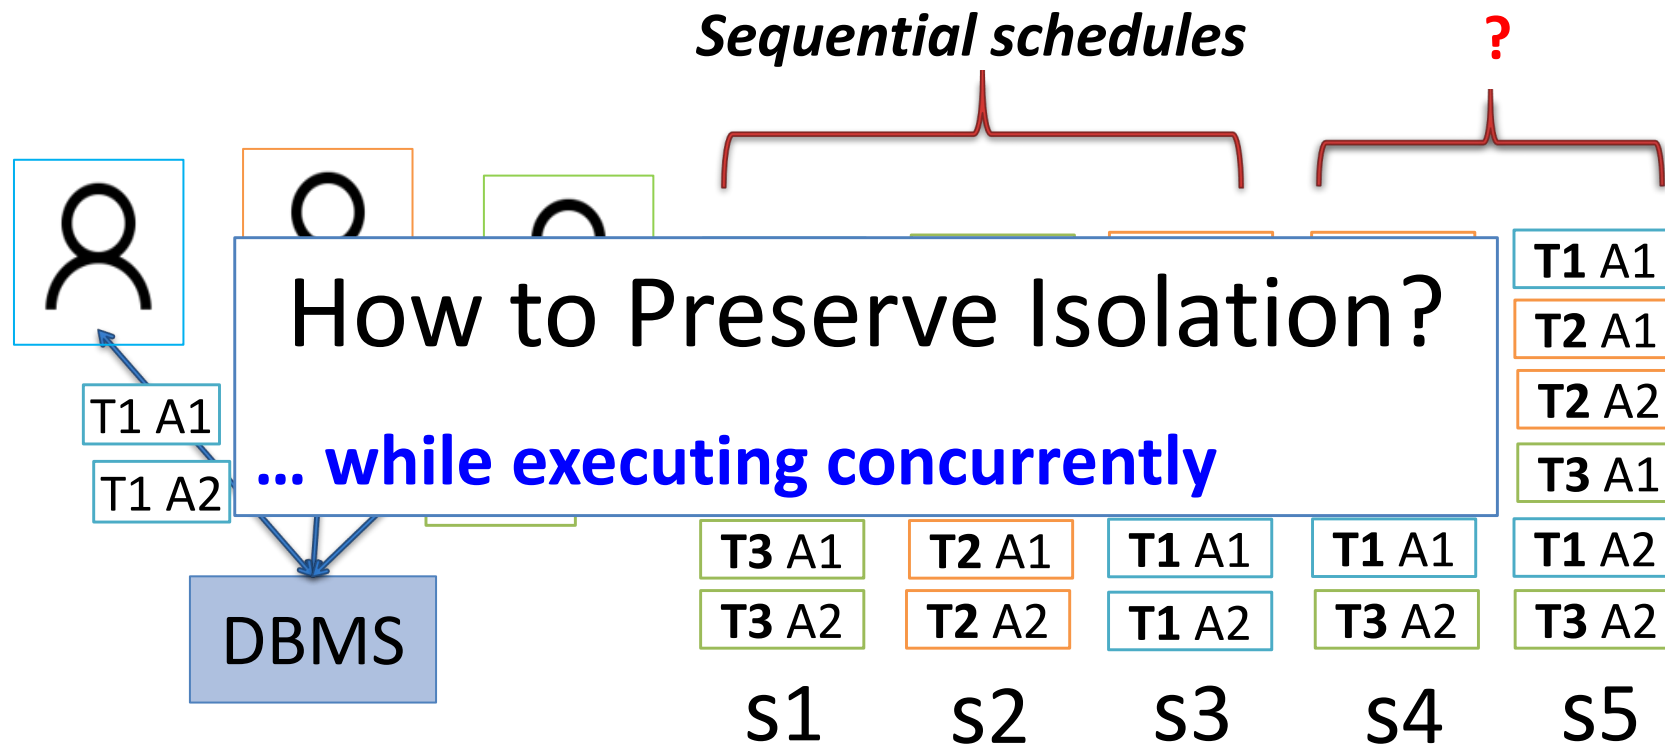
Operations may be interleaved,  
but execution must be  
equivalent to *some* sequential  
(serial) order of all transactions

How?

DBMS achieve isolation by  
enforcing **all schedules to be serializable**

# Implementing Isolation: Solution A

- Only allow sequential schedules



# Implementing Isolation: Solution B

1. Generate a schedule
2. Check if it is serializable

**How to check if a schedule is serializable?**

# Abstraction for computing Isolation

- How do we check if a parallel execution is correct?
  - It depends on the exact calculation of each program execution.
  - Too complicated\* to analyze the exact calculation that each calculation performs.
- Thus, we refer only to the abstraction of the programs:
  - Reading operations from the database.
  - Writing actions into the database.
  - The sequence, on the timeline, in which the actions are performed.

\*In fact, it is not possible

Programs	
Read(A)	
A:=A-100	
	Read(C)
Write(A)	
	C:=C-100
Read(B)	
	Write(C)
B:=B+100	
	Read(B)
	B:=B+100
	Write(B)
Write(B)	

# Program abstraction

Actual program	Abstraction
Read(A) A:=A-100	Read(A)
Write(A)	Read(C)
Read(B)	Read(B)
Write(C)	Write(C)
B:=B+100	Read(B)
Read(B)	Write(B)
B:=B+100	
Write(B)	Write(B)

**We focus only on the read and write operations**

# How to Check if a Schedule is Serializable?

Approach: check ordering constraints

Schedule #1	
T1	T2
Read(A)	
	Read(C)
Write(A)	
	Write(C)
	Read(B)
	Write(B)
Read(B)	
Write(B)	

Schedule #2	
T1	T2
Read(A)	
	Read(C)
Write(A)	
Read(B)	
	Write(C)
	Read(B)
	Write(B)
Write(B)	

# How to Check if a Schedule is Serializable?

Is it serializable?

Schedule #3		Schedule #2	
T1	T2	T1	T2
Read(A)		Read(A)	
	Read(C)		Read(C)
Write(A)		Write(A)	
<b>Read(B)</b>		<b>Read(B)</b>	
	Write(C)		Write(C)
	Read(B)		Read(B)
Write(B)			<b>Write(B)</b>
		Write(B)	




# How to check if schedule is serializable?

- Definition: Precedence graph
  - Every transaction is a vertex.
  - There is an edge from  $T_i$  to  $T_j$  if and only if one of the following conditions is met:
    - $T_i$  performs Read( $x$ ) before  $T_j$  performs Write ( $x$ )
    - $T_i$  performs Write( $x$ ) before  $T_j$  performs Read( $x$ )
    - $T_i$  performs Write ( $x$ ) before  $T_j$  performs Write( $x$ )
  - Note: read( $X$ ) read( $X$ ) do not add an edge



# How to check if schedule is serializable?

- Scheduling is serializable if there are no circles in the precedence graph.
  - How to check if there are circuits (algorithm course)?
  - Is it possible to have a serializable schedule in a situation where there are circuits?

# Precedence graph - illustration



Equivalent serial scheduling:



# Implementing Isolation: Solution C

- Goal: create serializable schedules

## Approach #1: using **locks**

- What are we locking and why?
- What types of locks do we need?
- When to use which lock?

# Exclusive and Shared Locks

- **Exclusive lock**
  - Only one transaction can hold it at time
  - When to use?
- **Shared lock**
  - Many transaction can hold this lock
  - When to use?

# When we lock a key?

- To read item A you must first lock it in a shared lock.
  - It is possible to release the lock after reading.
- To write item A you must first lock it with an exclusive lock.
  - It is possible to release the lock after writing.
- If after reading need to write, then transaction can upgrade a shared lock to an exclusive lock.
- Requests for locks are directed to the lock manager.

# When is it allowed to lock a key?

- Multiple transactions can simultaneously hold a shared lock on item A.
  - Provided there is no transactions holding an exclusive lock on A.
- Only one transaction can hold an exclusive lock on item A.
  - Then no other transaction has a shared or exclusive lock on A.
- The lock manager makes sure the rules are followed.

# Lock manager

- The lock manager handles lock requests and lock releases.
- Holds a table that has an entry for each item A with:
  - List of transactions currently holding a lock on A.
  - Type of the lock.
  - Pointer to queue of lock requests on A.
- Receiving a lock and releasing a lock should be done as atomic operations.
- A common lock can be upgraded to an exclusive lock.



# How and When to Lock?

## Schedule #2

T1	T2
Read(A)	
	Read(C)
Write(A)	
Read(B)	
	Write(C)
	Read(B)
	Write(B)
Write(B)	

Lock Type	Owner	Data item
Exclusive	T1	A
Exclusive	T2	C
Shared	T1	B
Exclusive	T2	B
Exclusive	T1	B

**T1 must release its shared key  
on B before T2**

**T2 must release its exclusive  
key on B before T1**

# Concurrency Control with Locking

- Goal: **create serializable schedules**
- Approach: using locks

## 2 Phase Locking (2PL)!

# Two-Phase Locking (2PL)

# Two-Phase Locking (2PL)

- Transaction must get:
  - Shared lock before reading an item.
  - Exclusive lock before writing an item.
- Once a transaction releases a lock it cannot request additional locks.
- That is why the protocol is called "two-phase locking"
  - First a phase of making locks.
  - Followed by a phase of release of locks.

# 2PL Example

Is 2PL?

T1	T2
<sup>S(A)</sup> Read(A)	
	<sup>S(C)</sup> Read(C)
<sup>X(A)</sup> Write(A)	
<sup>S(B)</sup> Read(B)	
	<sup>X(C)</sup> Write(C)
	<sup>S(B)</sup> Read(B)
	<sup>X(B)</sup> Write(B)
<sup>X(B)</sup> Write(B)	

Not a 2PL!

Is 2PL?

T1	T2
<sup>S(A)</sup> Read(A)	
	<sup>S(C)</sup> Read(C)
<sup>X(A)</sup> Write(A)	
<sup>S(B)</sup> Read(B)	
<sup>X(B)</sup> Write(B)	
T1 release all its keys	
	<sup>X(C)</sup> Write(C)
	<sup>S(B)</sup> Read(B)
	<sup>X(B)</sup> Write(B)

2PL

# The Downside of Using Locks Is ...

## Deadlocks

How handle deadlocks?

*Detect &  
handle*

*Avoid*

# Deadlock

- A state of deadlock occurs when there is a circle of transactions, so that every transaction in the circle waits for a lock that is currently held by the next transaction in the circle.
- Two ways to deal with deadlock:
  - Prevention of deadlock.
  - Detection of deadlock and cancellation.

# Is there a Deadlock here?

T1: S(A), R(A), S(B)

T2: X(B), W(B) X(C)

T3: S(C), R(C) X(A)

T4: X(B)




# How to find a Deadlock?

T1: S(A), R(A), S(B)

T2: X(B), W(B) X(C)

T3: S(C), R(C) X(A)

T4: X(B)



# Detection of deadlock – option 1

- Build a "wait for" graph
  - Node for every transaction.
  - from  $T_1$  to  $T_2$  if  $T_1$  is waiting for a lock that  $T_2$  currently holds.
- From time-to-time check if there is a circle in the graph (which means a state of deadlock)
- If there is a circle, cancel one of the transaction that is on the circle, and start it again.

# Detection of deadlock – option 2

- Timeout method - after a certain amount of time without progress in making the transactions, it can be assumed that there is a Deadlock.
- Release from Deadlock is done by performing ABORT on at least one of the transactions.