

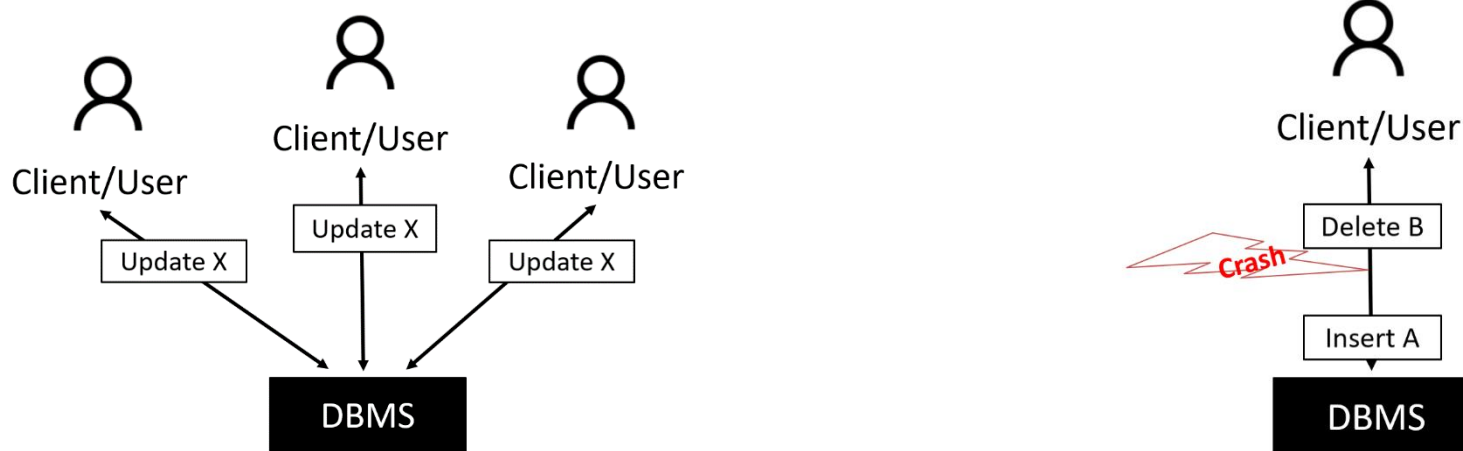
# Introduction to Operating Systems and SQL for Data Science

---

## Practice 12 – Transactions

# Challenges in a real database

Challenges in a DB with **many users** that may **fail**



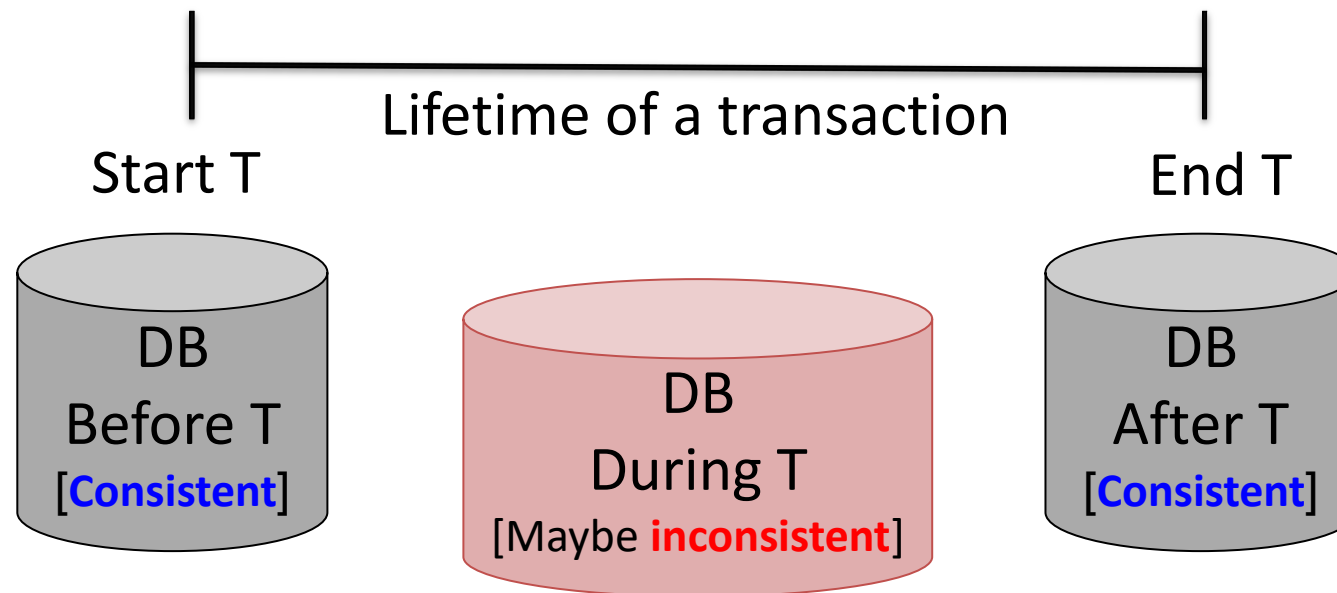
**Transactions to save the day!**

# Transactions

# Transactions to the Rescue!

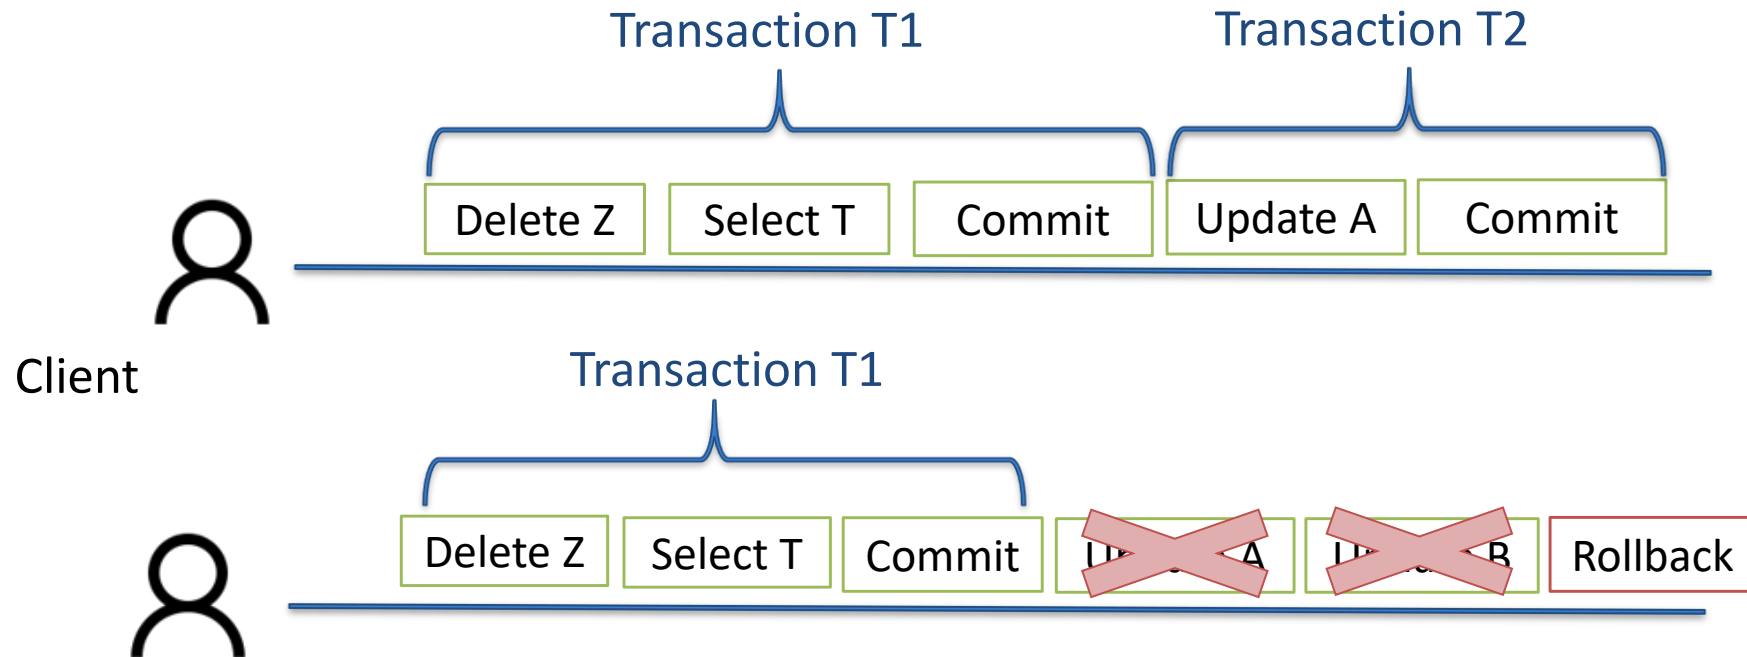
**Transaction:** a sequence of SQL operations that are considered as a single unit

- In a transaction either **all actions are done or none**
- Transitions the DB from one consistent state to another

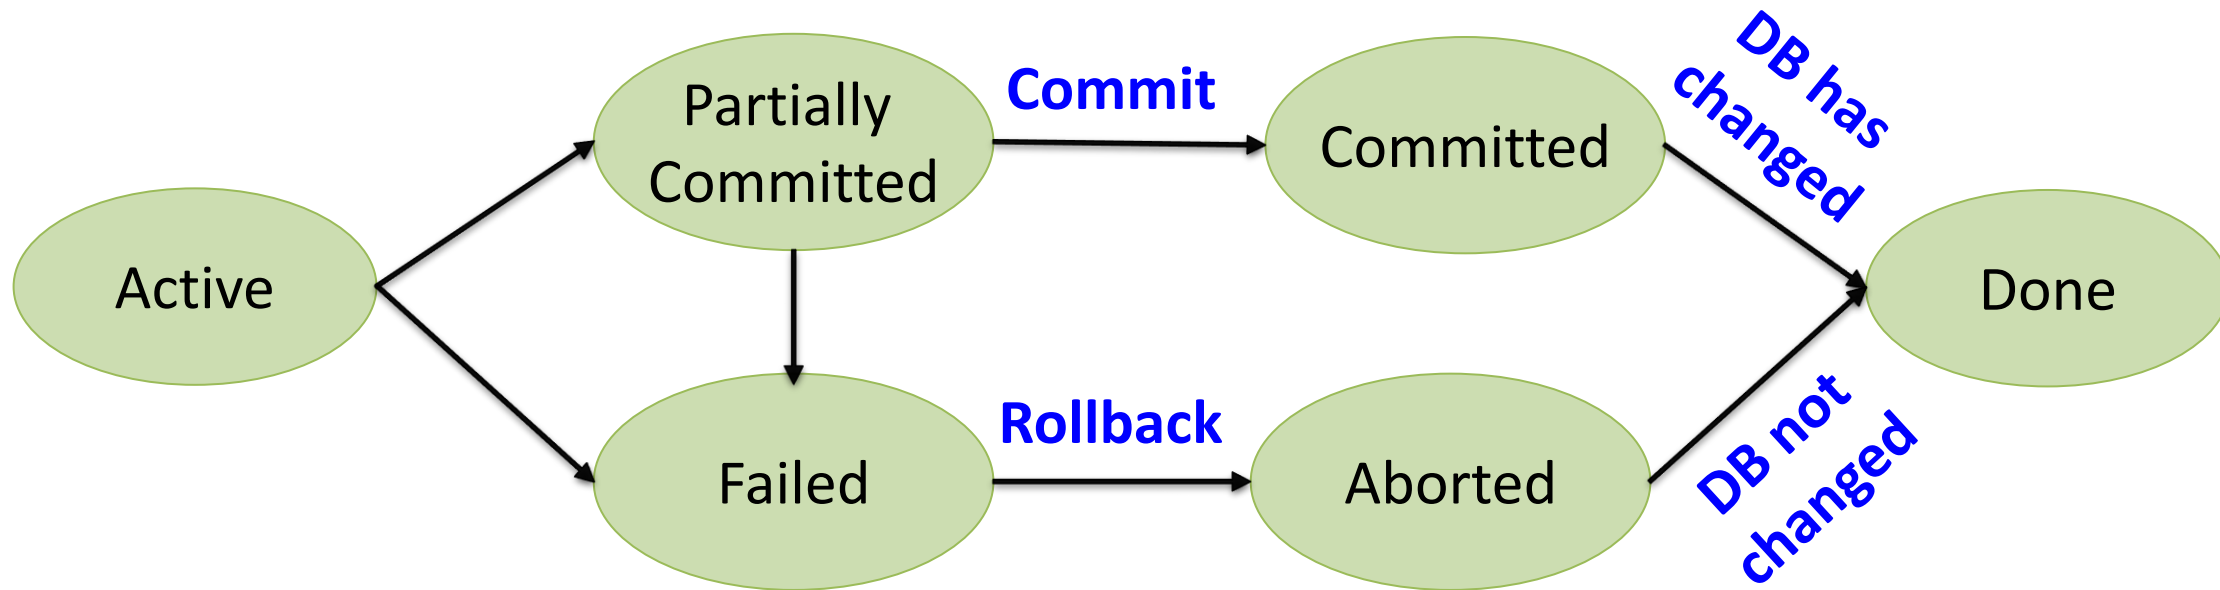


# Transaction mechanism

- How we create transactions?
  1. We apply SQL actions.
  2. When we want to “close” a Transaction we call this command **Commit**.
  3. If we want to “cancel” a started Transaction we do a **Rollback**.
- In many languages we can define auto commit



# Transaction lifecycle



# ACID: Desired Properties of Transactions

- Atomicity
- Consistency
- Isolation
- Durability

# Serializable schedule

- A schedule is a serializable schedule if it is equivalent to some serial schedule.
- Need a way to determine that a schedule is serializable without knowing what calculations are being performed.
- That is, for every possible calculation, the timing should be equivalent to a serial schedule.



# Equivalent schedules

## Two schedules are equivalent if:

- They are composed of the same actions.
- Both schedules have the same effect on the database, meaning that the database has the same values at the end of each of the schedules.
- Both schedules produce the same output to the user, i.e., present the user with the same values for each item they read.

# How to check if schedule is Conflict serializable?

- Definition: Precedence graph
  - Every transaction is a vertex.
  - There is an edge from  $T_i$  to  $T_j$  if and only if one of the following conditions is met:
    - $T_i$  performs Read(x) before  $T_j$  performs Write (x)
    - $T_i$  performs Write(x) before  $T_j$  performs Read(x)
    - $T_i$  performs Write (x) before  $T_j$  performs Write(x)
  - Note: read(X) read(X) do not add an edge

# How to check if schedule is conflict serializable?

- A Schedule is conflict serializable iff there are no circles in the precedence graph.

# Example

Given 4 transaction over X Y Z P

$T_1 = R_1(X) \ W_1(X) \ R_1(P) \ W_1(P)$

$T_2 = R_2(Y) \ W_2(Y) \ R_2(P)$

$T_3 = R_3(Z) \ R_3(Y) \ W_3(Z) \ R_3(P)$

$T_4 = R_4(X) \ W_4(X) \ R_4(Z)$

And the following schedule S:

$R_4(X)W_4(X)R_1(X)W_1(X)R_2(Y)W_2(Y)R_1(P) \dots$


$\dots W_1(P)R_2(P)R_3(Z)R_3(Y)W_3(Z)R_4(Z)R_3(P)$

**Build dependency graph of S**


$$R_4(X)W_4(X)R_1(X)W_1(X)R_2(Y)W_2(Y)R_1(P)W_1(P)R_2(P)R_3(Z)R_3(Y)W_3(Z)R_4(Z)R_3(P)$$

$R_4(X)W_4(X)R_1(X)W_1(X)R_2(Y)W_2(Y)R_1(P)W_1(P)R_2(P)R_3(Z)R_3(Y)W_3(Z)R_4(Z)R_3(P)$

$T_4$	$T_3$	$T_2$	$T_1$
READ X			
WRITE X			
			READ X
			WRITE X
		READ Y	
		WRITE Y	
			READ P
			WRITE P
		READ P	
	READ Z		
	READ Y		
	WRITE Z		
READ Z			
	READ P		



Is S conflict serializable?



Is S conflict serializable? Yes since there are no circles in the dependency graph



# concurrency control: locks

# Locks

**Locks** are data structures and protocols designed to prevent pre-creation of conflicts between transactions and prevent circuits in the conflict graph .

**X-LOCK (Exclusive Lock)**: Write lock.

Only one transaction can hold a writing lock on the same item.

**S-LOCK (Shared Lock)**: Read lock.

Several transactions can hold S-LOCK on the same item in parallel, as long as the item does not have a write lock (X-LOCK)

# Access and lock conditions

- A transaction can write item A only after locking x-lock on A
- A transaction can read item A only after locking s-lock on A
- A transaction cannot x-lock A if there is already **any** lock on it
- A transaction cannot s-lock A if there is already x-lock on it
- Each transaction must release its locks before it finishes.

# Two-Phase Locking (2PL)

# Two-Phase Locking (2PL)

- A transaction holds 2PL requirements if:
  - It satisfied access and lock conditions
  - All locking actions happens **before** releases actions (once the transaction released a lock it cannot lock any longer)
- That is why the protocol is called "two-phase locking"
  - First a phase of making locks.
  - Followed by a phase of release of locks.

# Two-Phase Locking (2PL)

- When all transactions in a schedule  $S$  are 2PL we say that the scheduling is 2PL.

# Two-Phase Locking (2PL)

- When all transactions in a schedule  $S$  are 2PL we say that the scheduling is 2PL.

# An example:





# The Downside of Using Locks Is ...

## Deadlocks

How handle deadlocks?

*Detect &  
handle*

*Avoid*

# Deadlock

- A state of deadlock occurs when there is a circle of transactions, so that every transaction in the circle waits for a lock that is currently held by the next transaction in the circle.
- Two ways to deal with deadlock:
  - Prevention of deadlock.
  - Detection of deadlock and cancellation.


# Deadlock detection

- We will draw a graph:
  - Transaction  $T_i$  as vertexes
  - Edge  $T_i \rightarrow T_j$  if  $T_i$  waits for a lock held by  $T_j$

There is a deadlock iff there is a circle in the graph

# Deadlock detection

T1	T2	T3
X-LOCK A		
	X-LOCK B	
		X-LOCK C
S-LOCK C		
	X-LOCK A	
		S-LOCK B



# Deadlock detection

Given the following 2 transactions:

T1: R(A), R(B), W(B)

T2: R(B), R(A), W(A)

- a. Add lock and release actions so each transaction satisfied 2PL.
- b. There is a schedule S with T1 and T2 that ends up with a deadlock?

Even if both transactions are 2PL. If so – give an example, if not- explain

Given the following 2 transactions:

T1: R(A), R(B), W(B)

T2: R(B), R(A), W(A)

a. Add lock and release actions so each transaction satisfied 2PL.

Given the following 2 transactions:

T1: R(A), R(B), W(B)

T2: R(B), R(A), W(A)

a. Add lock and release actions so each transaction satisfied 2PL.

T1	T2
S-LOCK (A)	S-LOCK (B)
R (A)	R (B)
X-LOCK (B)	X-LOCK (A)
R (B)	R (A)
W (B)	W (A)
UNLOCK (A)	UNLOCK (B)
UNLOCK (B)	UNLOCK (A)

b. There is a schedule S with T1 and T2 that ends up with a deadlock?

Even if both transactions are 2PL. If so – give an example, if not- explain

Yes. Look at the following schedule:

T1	T2
S-LOCK (A)	
	S-LOCK (B)
	R (B)
R (A)	
X-LOCK (B)	
	X-LOCK (A)
	R (A)
	W (A)
R (B)	
W (B)	
UNLOCK (A)	
UNLOCK (B)	
	UNLOCK (B)
	UNLOCK (A)



There is a circle in the graph => there is a deadlock